



Assignment Report on:

“Architectural Design for Facebook”

Submitted to:

**Mridha Md. Nafis Fuad
Lecturer
Institute of Information Technology
University of Dhaka**

Submitted by:

**Pal Swadhin (BSSE 1302)
Mostafizur Rahman (BSSE 1320)
Rony Majumder (BSSE 1325)
Ibne Bin Rafid (BSSE 1330)
Fareya Azam (BSSE 1331)**

**Course- Software Design (SE-606)
Institute of Information Technology**

Submitted on 16th January, 2025

Table of Contents

Architectural Design for Facebook: microservices and distributed designs.....	4
1. High-Level Modules.....	4
2. Key Microservices.....	4
2.1 User Management.....	4
2.2 Content Management.....	4
2.3 Social Graph Management.....	4
2.4 Feed Management.....	5
2.5 Notification System.....	5
2.6 Messaging.....	5
2.7 Search.....	5
2.8 Ads System.....	5
2.9 Analytics.....	5
2.10 Admin Tools.....	5
3. Distributed Systems Design.....	6
3.1 Databases.....	6
3.2 Storage.....	6
3.3 Message Queue.....	6
3.4 Real-Time Systems.....	6
3.5 Load Balancing.....	6
3.6 Caching.....	6
4. Component Interaction.....	7
4.1 Frontend.....	7
4.2 Backend.....	7
4.3 Microservices Communication.....	7
5. Deployment.....	7
5.1 Containerization.....	7
5.2 Monitoring and Observability.....	7
Connected Micro Services.....	8
Replication and Synchronization Strategies.....	11
1. User Database.....	11
9. Friendship DB.....	12
2 & 3 . Notification DB & Feed Cache DB.....	13
Notification DB.....	13
Feed Cache DB.....	13
4. Media Database.....	14
6. React Content DB.....	15
5. Post Content DB.....	16

7. Comment Content DB.....	16
8. Search Index DB.....	16
10. Ads System DB.....	18
Global Load Balancer.....	19
Facebook Microservices Feature Collaboration.....	23
Communication Technologies Used.....	23
1. Synchronous Communication.....	23
2. Asynchronous Communication.....	23
3. Data Synchronization.....	23
Feature-by-Feature Demonstration.....	25
1. User Authentication and Profile Management.....	25
2. Post Creation and Timeline Update.....	26
3. Social Graph Operations (Friend Requests/Following).....	27
4. Feed Generation and Content Ranking.....	27
5. Messaging and Real-Time Communication.....	28
6. Ad System Integration.....	29
7. Search Functionality.....	29
Cross-Cutting Concerns.....	30
Data Consistency.....	30
Security.....	30
Monitoring and Logging.....	30
Error Handling.....	30
Performance Optimization.....	30

Architectural Design for Facebook: microservices and distributed designs

1. High-Level Modules

- **User Management:** Handles registration, authentication, and profile management.
- **Content Management:** Deals with posts, photos, videos, comments, and likes.
- **Social Graph Management:** Manages user connections (friends, followers).
- **Feed Management:** Personalized news feed generation.
- **Notification System:** Real-time alerts for likes, comments, friend requests, etc.
- **Messaging:** Instant messaging and group chat.
- **Search:** Global search for users, pages, groups, and posts.
- **Ads System:** Ad creation, targeting, and analytics.
- **Analytics:** Insights for user behavior and content performance.
- **Admin Tools:** Moderation and compliance enforcement.

2. Key Microservices

Each module above can be broken into smaller microservices:

2.1 User Management

- **Authentication Service:** OAuth, session management, password recovery.
- **User Profile Service:** CRUD operations for user details.
- **Privacy Settings Service:** Controls user permissions and privacy settings.

2.2 Content Management

- **Post Service:** Create, read, update, and delete (CRUD) posts.
- **Media Service:** Handles uploads and storage of images and videos.
- **Comment Service:** Manages comments on posts.
- **Reaction Service:** Like, love, or any reaction type.

2.3 Social Graph Management

- **Friendship Service:** Manage friend requests and relationships.
- **Follow Service:** Handles follower-following relationships.

2.4 Feed Management

- **Feed Aggregation Service:** Collects posts for feeds.
- **Recommendation Engine:** Suggests content based on user behavior.
- **Ranking Service:** Prioritizes posts in feeds.

2.5 Notification System

- **Push Notification Service:** Sends notifications to devices.
- **Inbox Notification Service:** Stores notifications in user accounts.

2.6 Messaging

- **Real-time Chat Service:** WebSocket-based instant messaging.
- **Group Messaging Service:** Manages multi-user chat.

2.7 Search

- **Search Indexing Service:** Builds indexes for fast searching.
- **Search Query Service:** Handles user queries.

2.8 Ads System

- **Ad Targeting Service:** Matches ads to user demographics.
- **Billing Service:** Manages advertiser payments.

2.9 Analytics

- **Event Logging Service:** Tracks user interactions.
- **Behavior Analytics Service:** Analyzes user data.

2.10 Admin Tools

- **Moderation Service:** Flags inappropriate content.
- **Reporting Service:** Manages user-reported issues.

3. Distributed Systems Design

3.1 Databases

- **User Database:** Stores user profiles (e.g., PostgreSQL, MySQL).
- **Content Database:** Stores posts, comments, and media metadata.
- **Social Graph Database:** Graph database (e.g., Neo4j) for relationships.
- **Feed Cache:** Redis or Memcached for feed generation.
- **Search Index:** Elasticsearch for full-text search.

3.2 Storage

- **Media Storage:** S3-compatible storage (e.g., MinIO) for photos and videos.
- **Backup System:** Ensures data reliability.

3.3 Message Queue

- **Kafka:** For asynchronous tasks like feed updates and notifications.

3.4 Real-Time Systems

- **WebSockets:** For live notifications and messaging.
- **Pub/Sub:** Real-time feed updates (e.g., Google Pub/Sub).

3.5 Load Balancing

- **Reverse Proxy:** Nginx or HAProxy for distributing requests.
- **Global Load Balancer:** Ensures availability across regions.

3.6 Caching

- **Edge Caching:** CDNs for faster content delivery.
- **Application Caching:** Redis for reducing database hits.

4. Component Interaction

4.1 Frontend

- React, Vue, or Angular for UI.
- GraphQL or REST APIs to interact with the backend.

4.2 Backend

- Node.js, Python (Flask/Django), or Java (Spring Boot) for services.
- APIs exposed for microservices communication.

4.3 Microservices Communication

- **Internal APIs:** REST or gRPC.
- **Event-driven Communication:** Using Kafka or RabbitMQ.

5. Deployment

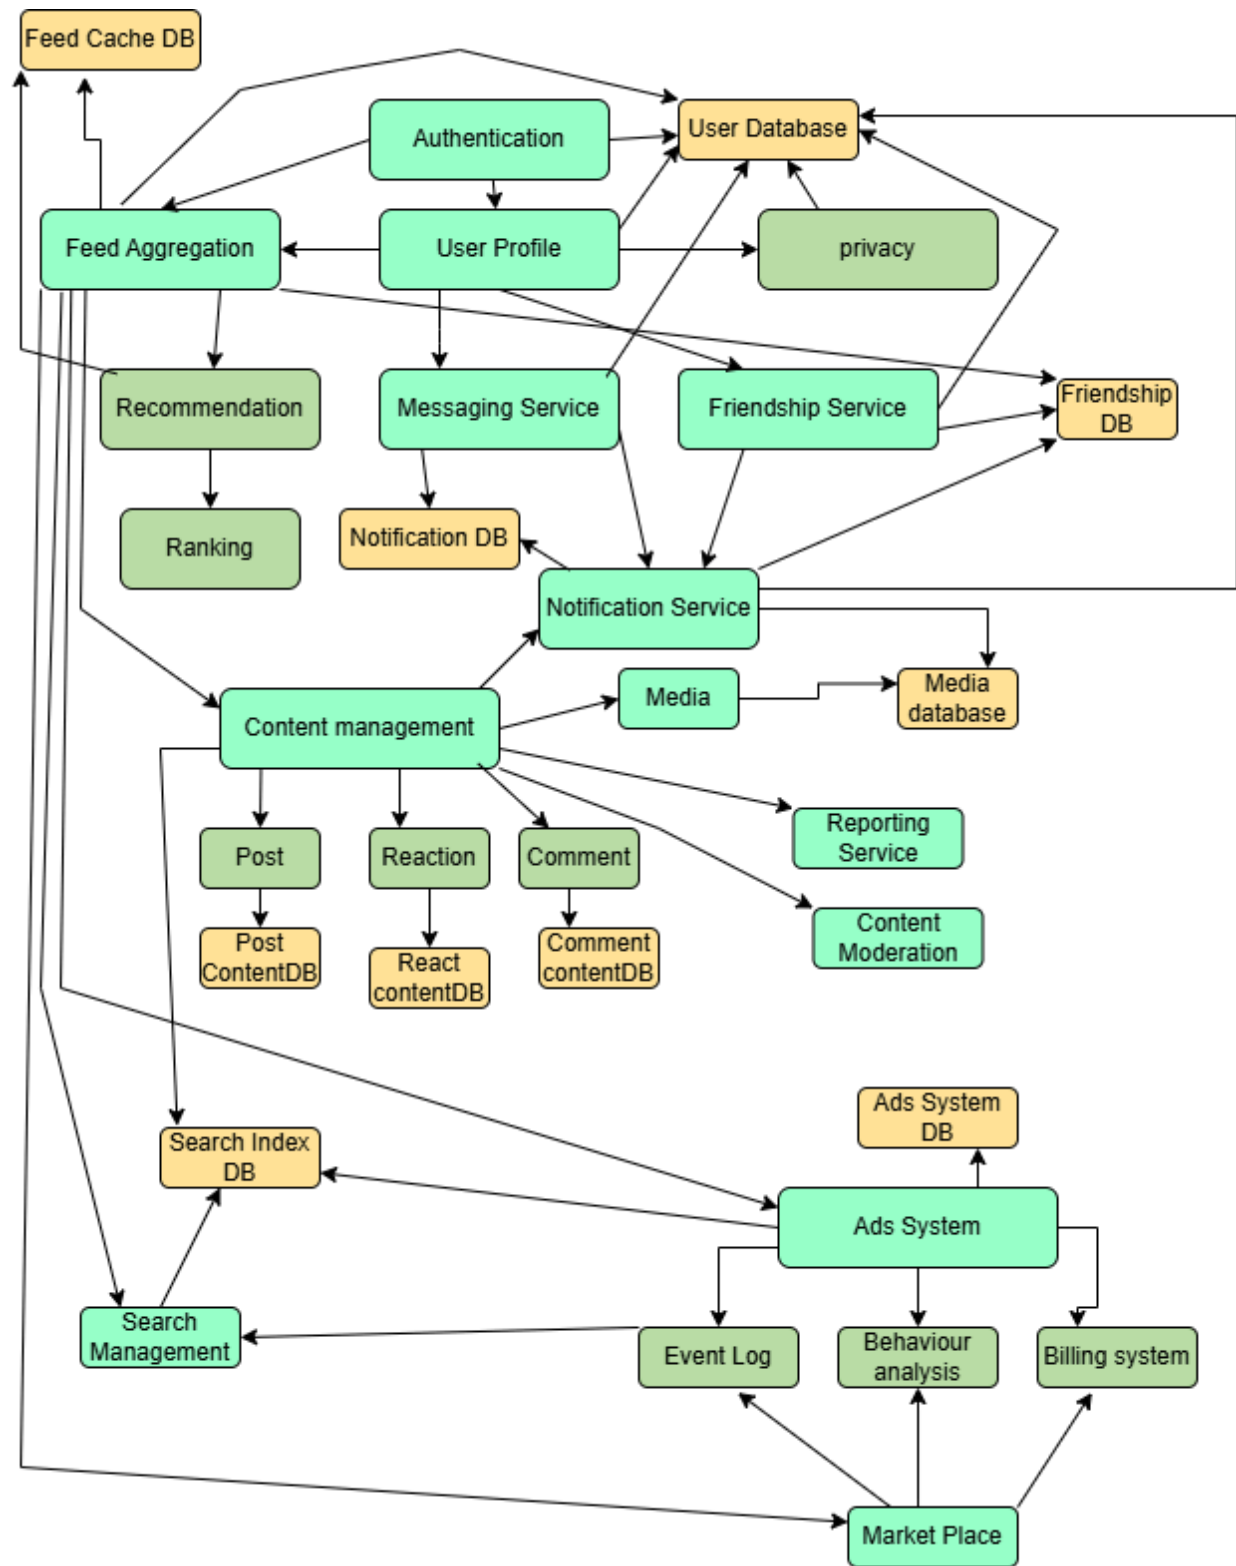
5.1 Containerization

- Docker containers for microservices.
- Kubernetes for orchestration.

5.2 Monitoring and Observability

- Prometheus and Grafana for metrics.
- ELK Stack for logging.

Connected Micro Services



Several databases are shared across multiple components in the system architecture, requiring efficient replication strategies to ensure data consistency and availability. For instance:

- The **User Database**, used by Authentication, User Profile, Privacy, and Friendship Service, leverages **Event-Driven Data Synchronization** to propagate user-related updates in near real-time while maintaining loose coupling between services.
- Content-related databases, such as the **Post Content DB**, **React Content DB**, and **Comment Content DB**, utilize **Change Data Capture (CDC)** to stream updates to dependent services like Feed Aggregation, ensuring the timely synchronization of posts, reactions, and comments.
- The **Media Database**, a shared resource, employs **API-Based Data Sharing** to avoid redundant storage and provide centralized access for services like Media and Reporting.
- To support low-latency operations, the **Feed Cache DB** and **Friendship DB** also use **Event-Driven Data Synchronization**, keeping feeds and notifications aligned with user interactions.
- For high-demand operations like those involving the **Ads System DB**, **Distributed Databases** ensure fault tolerance, scalability, and consistency across Ads-related components.

These replication strategies deliver a robust, scalable, and high-performance system that maintains consistency across its distributed microservices architecture.

Database and Components Table

#	Database	Connected Components	Replication Strategy	Reason
1	User Database	Authentication, User Profile, Privacy, Friendship Service	Event-Driven Data Synchronization	Decouples services and ensures eventual consistency for user-related data changes.
2	Feed Cache DB	Feed Aggregation, Recommendation, Ranking	Event-Driven Data Synchronization	Provides low-latency feed updates using cached data.
3	Notification DB	Notification Service, Messaging Service	Event-Driven Data Synchronization	Enables real-time notifications and message updates.
4	Media Database	Media Service, Reporting Service	API-Based Data Sharing	Avoids duplication of large files while maintaining centralized access for connected services.
5	Post Content DB	Content Management (Post Service)	Change Data Capture (CDC)	Streams updates to dependent services for real-time synchronization of posts.
6	React Content DB	Content Management (Reaction Service)	Change Data Capture (CDC)	Ensures reaction updates are propagated to related components like Feed and Notifications.
7	Comment Content DB	Content Management (Comment Service)	Change Data Capture (CDC)	Keeps comment-related data synchronized with Feed Cache and Search Index DB.
8	Search Index DB	Search Management, Content Management	Change Data Capture (CDC)	Maintains accurate and up-to-date search results for users and content.
9	Friendship DB	Friendship Service, Feed Aggregation, Notification Service	Event-Driven Data Synchronization	Ensures updates to friendships propagate to feeds and notifications for a seamless user experience.
10	Ads System DB	Ads System, Behavior Analysis, Event Log, Marketplace	Distributed Databases	Handles high data volumes and ensures consistency across Ads-related services for scalability.

Figure: The figure of how the databases, components, replication strategies, and reasons align with different microservices.

Replication and Synchronization Strategies

1. User Database

- **Connected Components:** Authentication, User Profile, Privacy, Friendship Service
- **Replication Strategy: Event-Driven Data Synchronization**
 - Changes to user data (e.g., profile updates) are published as events.
 - Services like Friendship or Feed Aggregation subscribe to user-related events for synchronization.

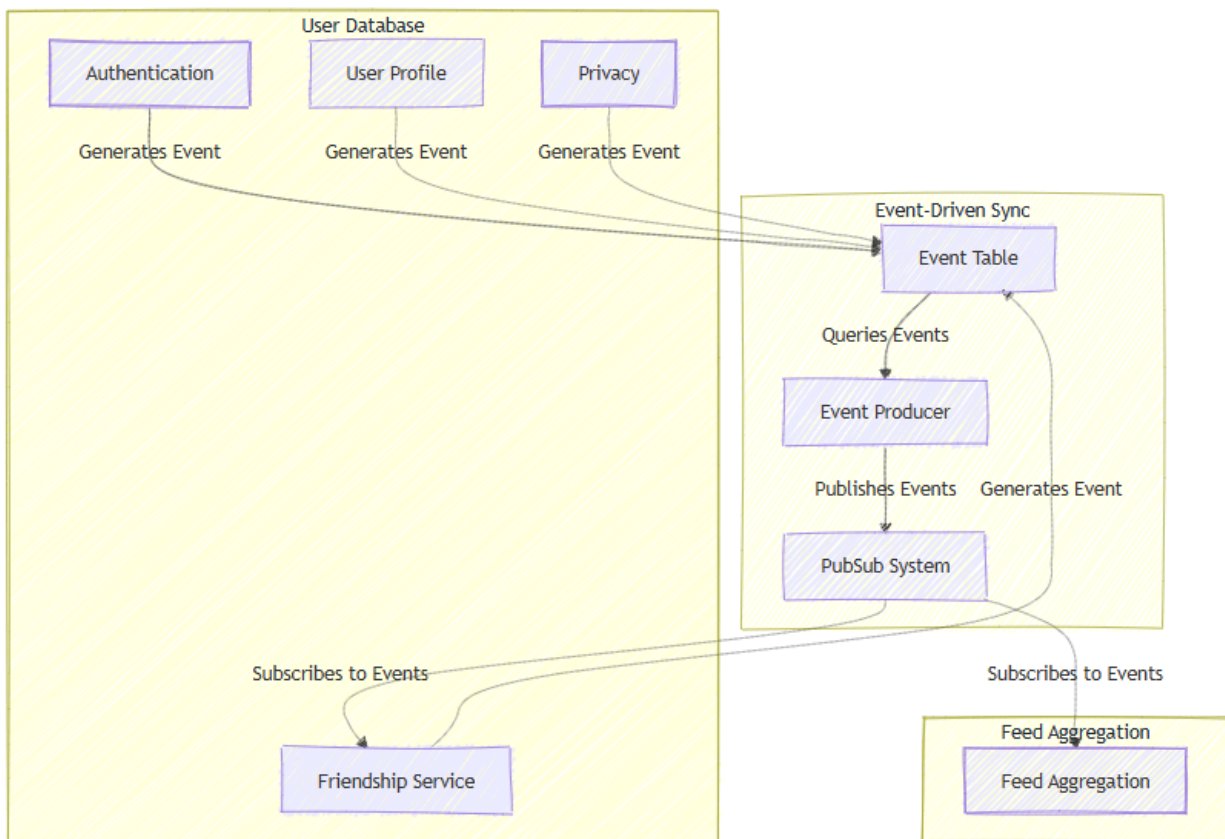


Figure: User Database Replication and Synchronization using event Event-Driven Data Synchronization

9. Friendship DB

- **Connected Components:** Friendship Service, Feed Aggregation, Notification Service
- **Replication Strategy: Event-Driven Data Synchronization**
 - Friendship updates (e.g., new connections) are published to Feed Aggregation and Notification systems.

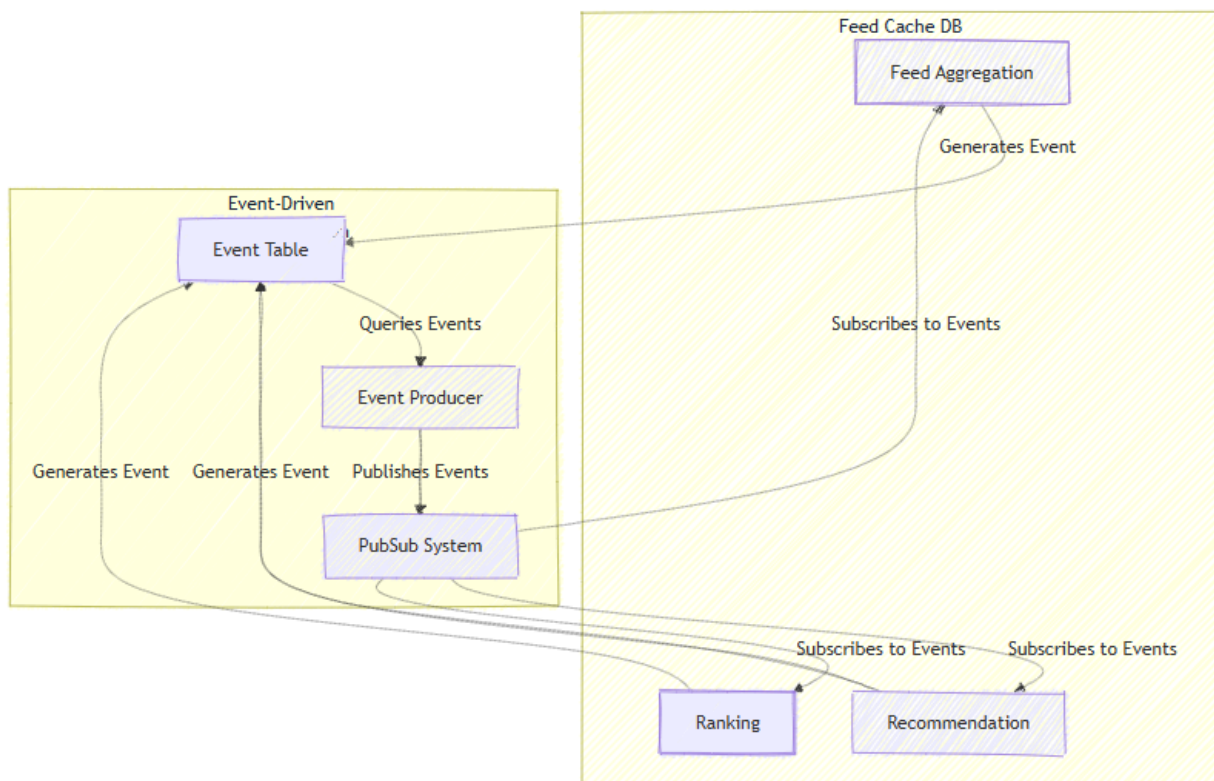


Figure: Friendship DB Replication and Synchronization using event Event-Driven Data Synchronization

2 & 3 . Notification DB & Feed Cache DB

Notification DB

- **Connected Components:** Notification Service, Messaging Service
- **Replication Strategy: Event-Driven Data Synchronization**
 - Notifications (e.g., new messages, friend requests) are generated via events.
 - Services like Messaging update their local data stores asynchronously.

Feed Cache DB

- **Connected Components:** Feed Aggregation, Recommendation, Ranking
- **Replication Strategy: Event-Driven Data Synchronization**
 - Feed changes due to new posts, reactions, or comments are processed asynchronously.
 - The cache is updated in near real-time for low-latency feed delivery.

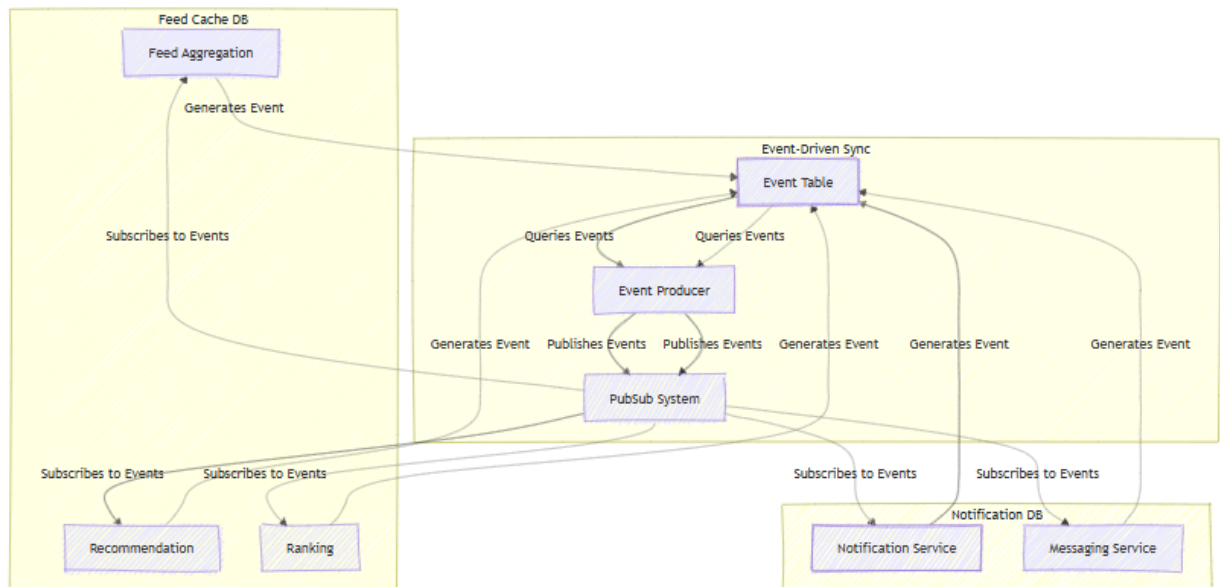


Figure: Notification DB & Friendship DB Replication and Synchronization using event Event-Driven Data Synchronization

4. Media Database

- **Connected Components:** Media Service, Reporting Service
- **Replication Strategy: API-Based Data Sharing**
 - Media is stored centrally, and metadata is shared via APIs.
 - This avoids large file duplication while maintaining access to connected services.

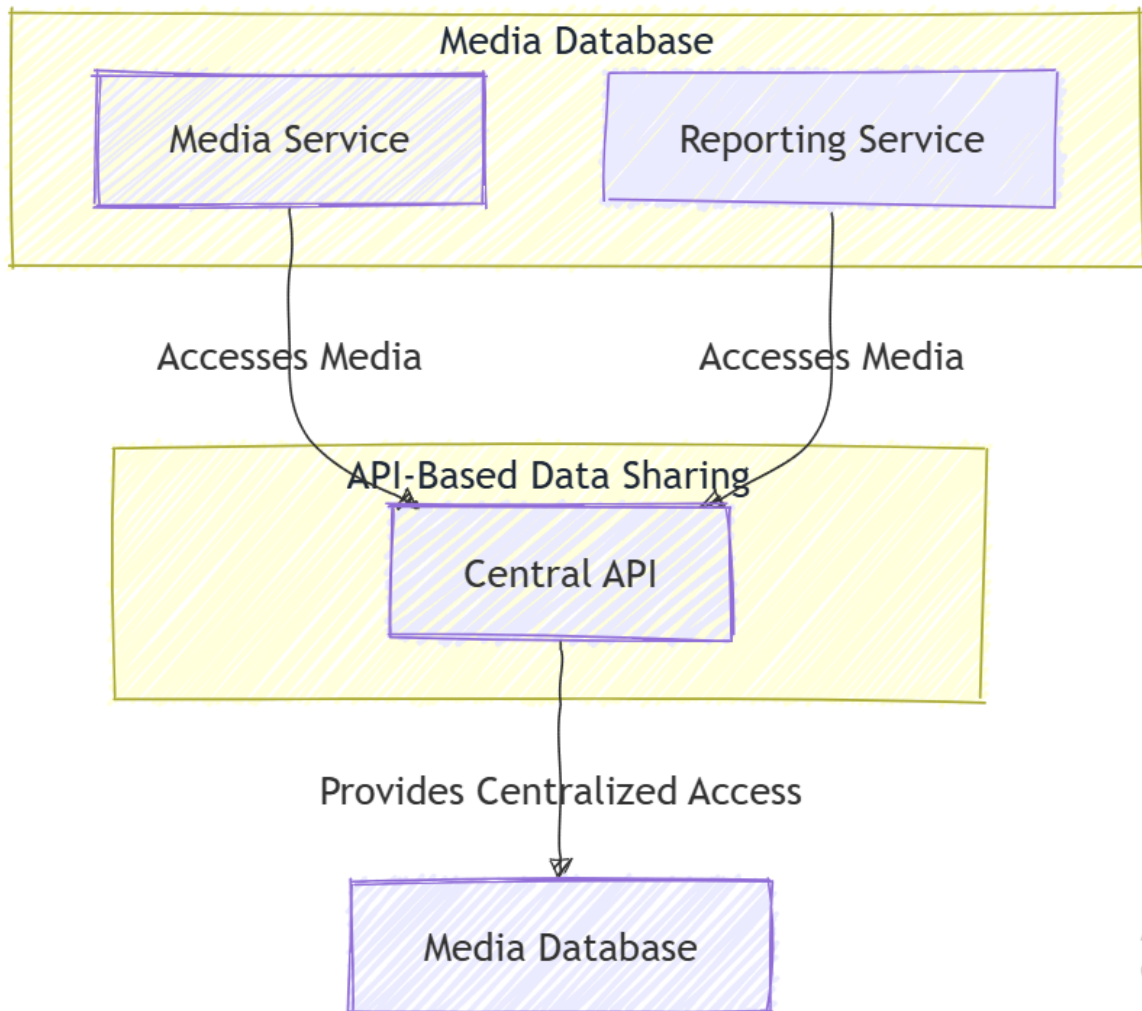


Figure: Media Database Replication and Synchronization using API-Based Data Sharing

6. React Content DB

- **Connected Components:** Content Management (Reaction Service)
- **Replication Strategy: Change Data Capture (CDC)**
 - Reaction updates are captured and synchronized with Feed Cache DB and Notification DB.

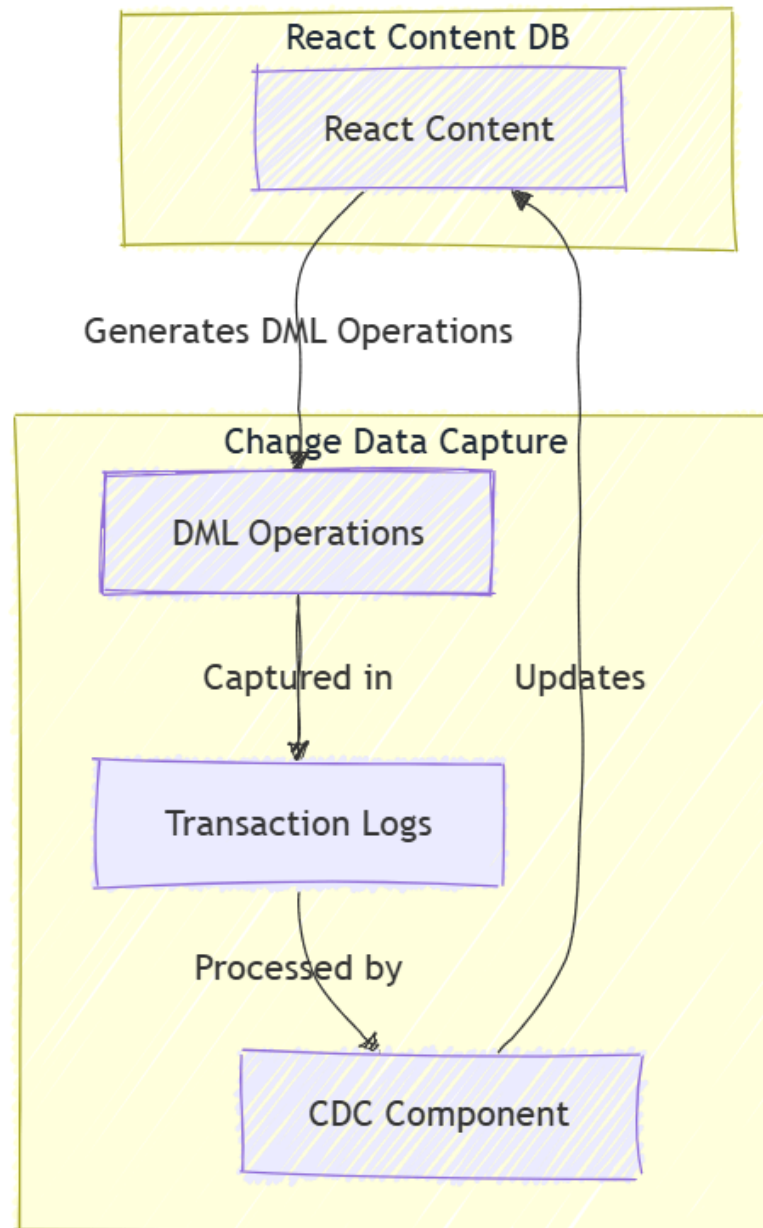


Figure: React Content DB Replication and Synchronization using Change Data Capture (CDC)

5. Post Content DB

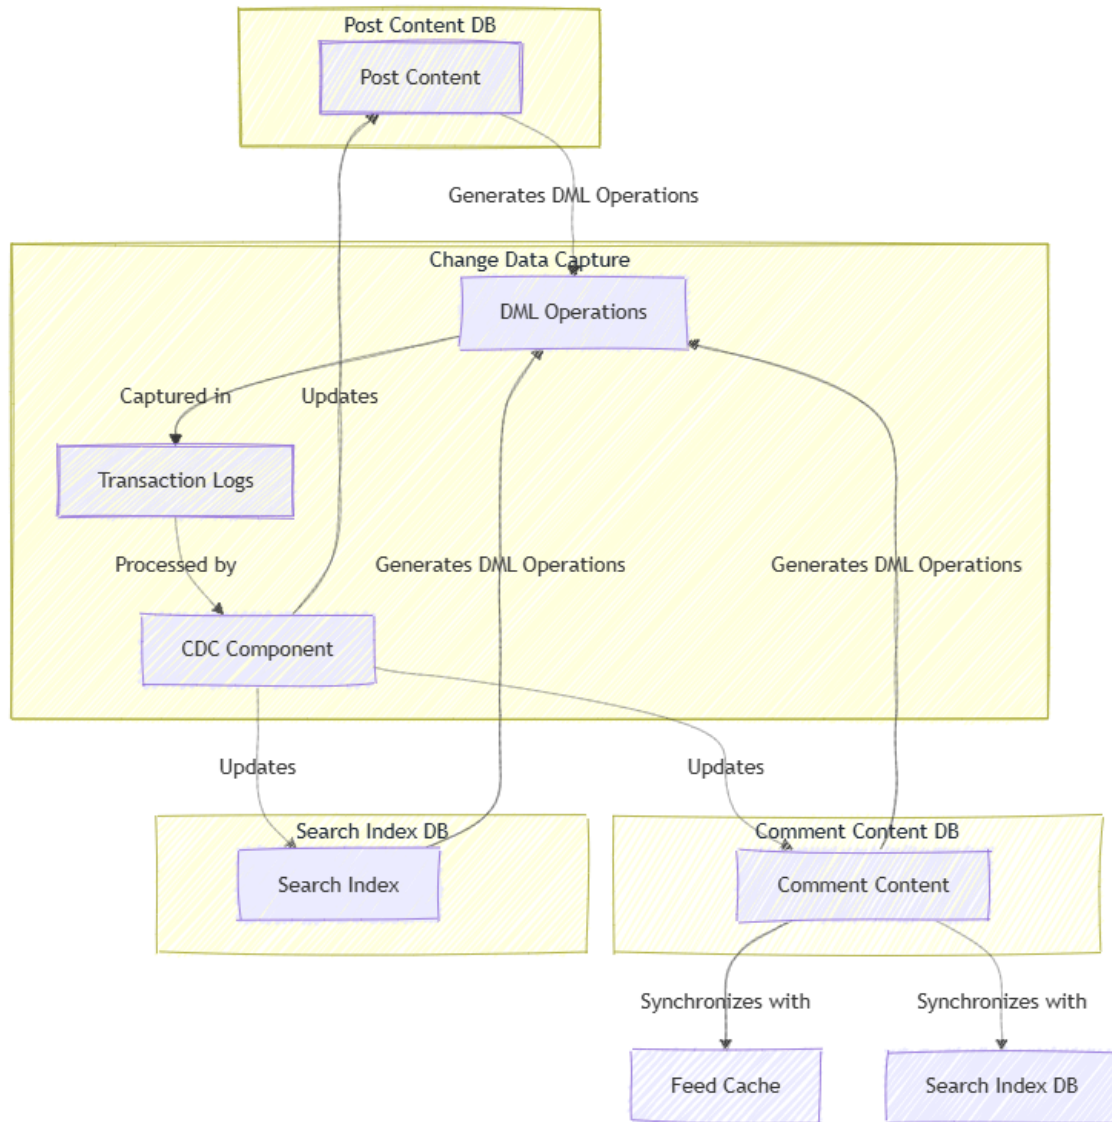
- **Connected Components:** Content Management (Post Service)
- **Replication Strategy: Change Data Capture (CDC)**
 - Changes to posts are streamed to Search Index DB and Feed Cache DB.
 - Tools like Debezium ensure updates are propagated to dependent systems.

7. Comment Content DB

- **Connected Components:** Content Management (Comment Service)
- **Replication Strategy: Change Data Capture (CDC)**
 - Similar to the React Content DB, comment changes propagate to dependent services.

8. Search Index DB

- **Connected Components:** Search Management, Content Management
- **Replication Strategy: Change Data Capture (CDC)**
 - Full-text search indexes are updated based on content changes (posts, comments, reactions).



*Figure: **Post Content DB**, **Comment Content DB** & **Search Index DB** Replication and Synchronization using Change Data Capture (CDC)*

10. Ads System DB

- **Connected Components:** Ads System, Behavior Analysis, Event Log, Marketplace
- **Replication Strategy: Distributed Databases**
 - A distributed database like Cassandra ensures scalability and low-latency access across multiple services.

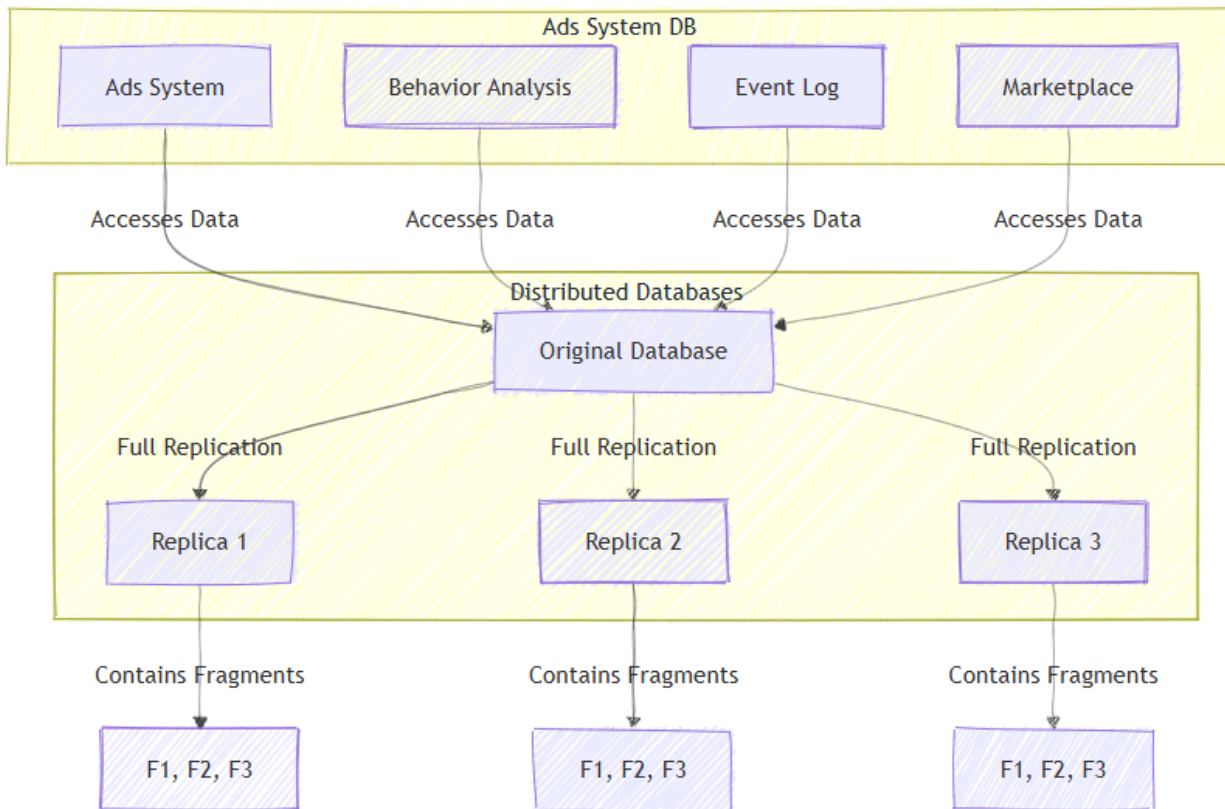


Figure: Ads System DB Replication and Synchronization using Distributed Databases

Global Load Balancer

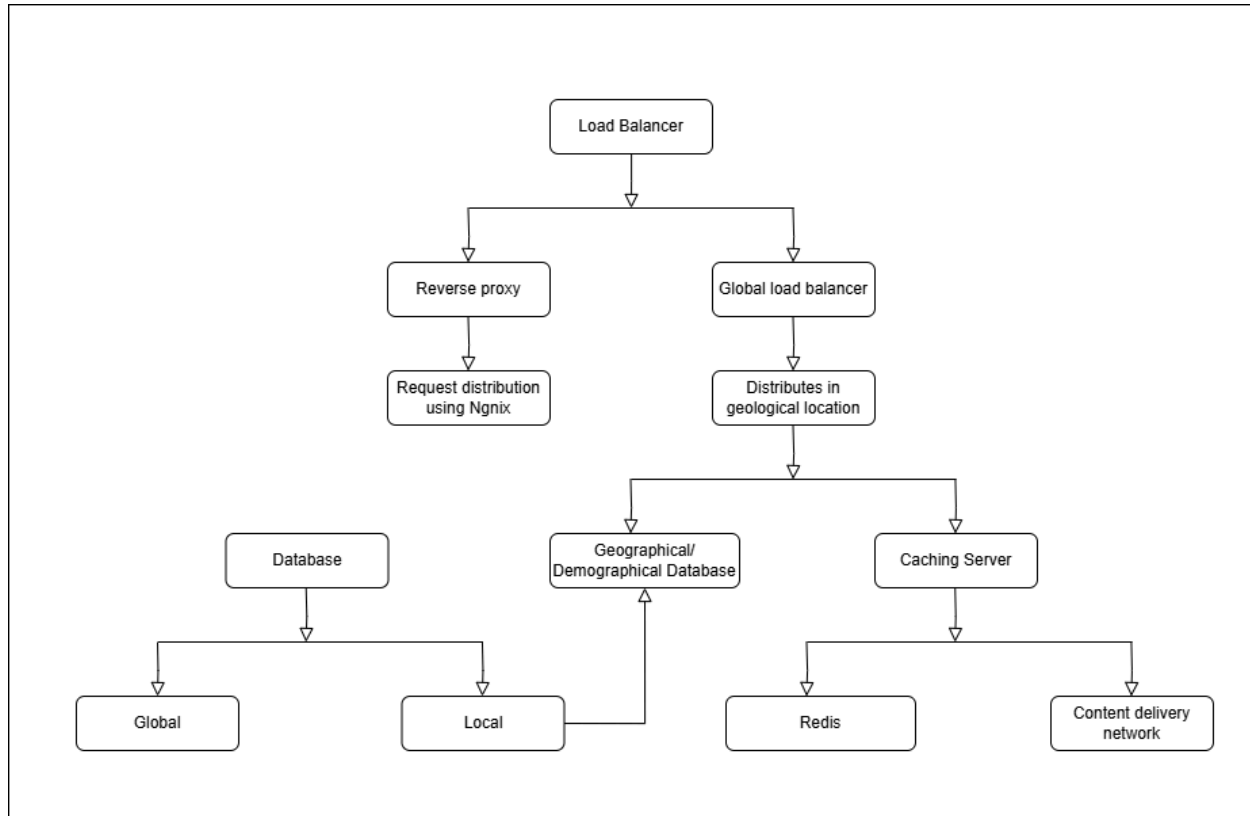


Fig: Load Balancer and Traffic Management System

This diagram illustrates a multi-layered architecture for handling web traffic efficiently:

1. **Load Balancer:** The primary component that manages incoming traffic.
2. **Reverse Proxy:** Uses Nginx to distribute requests to local databases and services.
3. **Global Load Balancer:** Distributes traffic based on geographical location.
4. **Geographical/Demographic Database:** Differentiates between global and local databases for efficient routing.
5. **Caching Server:** Uses Redis and Content Delivery Network (CDN) for faster data retrieval and reduced latency.

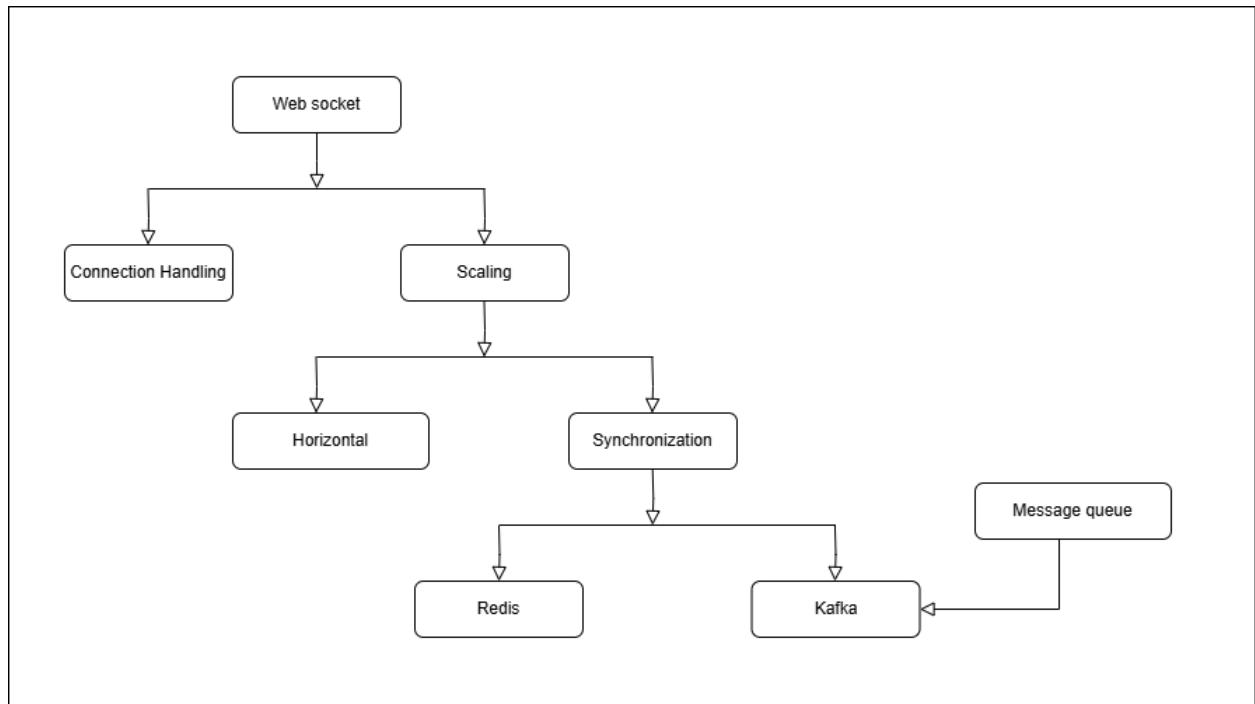
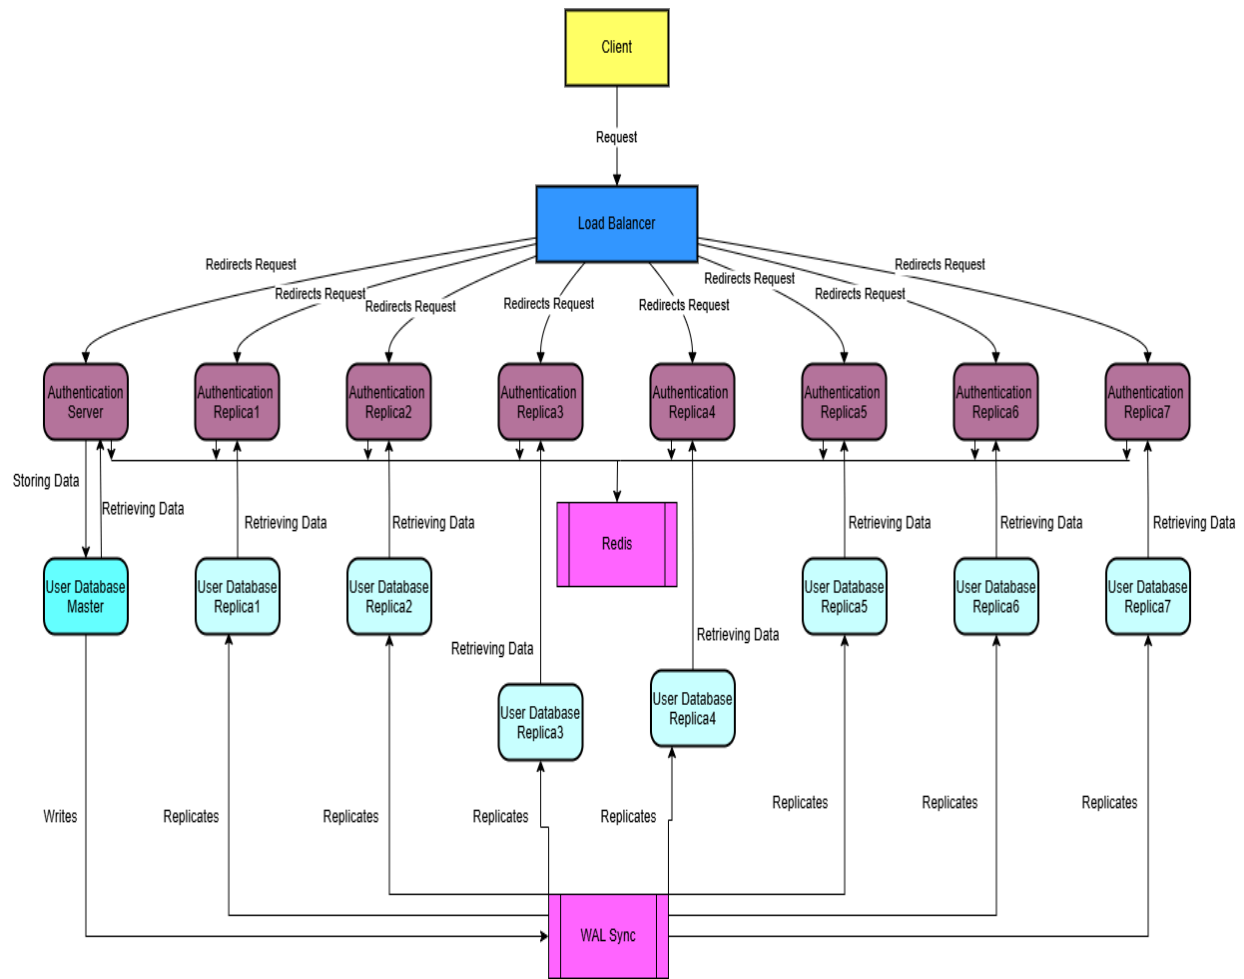


Fig: WebSocket Connection, Scaling, and Synchronization Workflow

This diagram illustrates a WebSocket-based architecture for handling real-time communication and scaling efficiently:

1. **WebSocket:** The core component managing real-time communication.
2. **Connection Handling:** Manages WebSocket connections from clients.
3. **Scaling:** Ensures the system can handle increasing demand:
 - **Horizontal Scaling:** Adds more servers or instances to distribute the load.
 - **Synchronization:** Maintains consistent data across servers.
4. **Redis:** Used for caching or session storage during synchronization.
5. **Kafka:** Used for message streaming and communication across distributed systems.
6. **Message Queue:** Handles asynchronous message processing to ensure reliability.



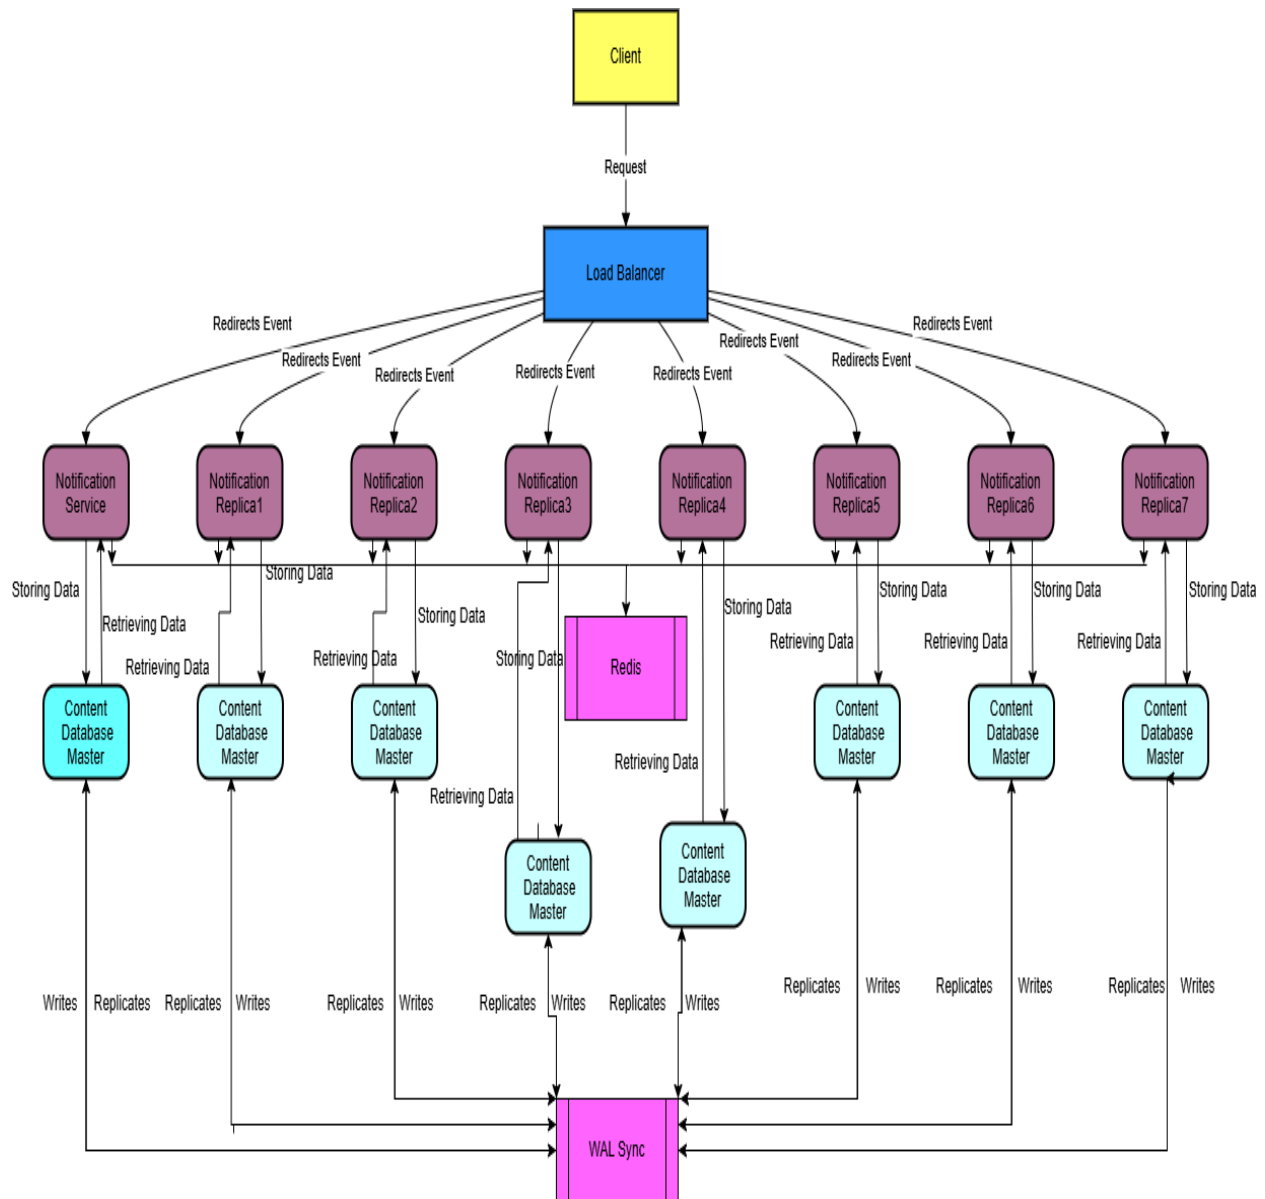
<https://drive.google.com/file/d/1k3eOdMCjpWq2ueZpsO8HO5Hxz1TWOs2c/view?usp=sharing>

Fig: Distributed Authentication and Database Replication System

This figure represents a distributed architecture for handling authentication requests and database synchronization with high scalability and reliability:

1. **Client:** Sends requests to the system.
2. **Load Balancer:** Distributes incoming requests across multiple authentication servers or replicas.
3. **Authentication Servers:**
 - A central authentication server manages requests.
 - Multiple replicas ensure scalability and fault tolerance, retrieving data from the user database replicas.
4. **User Database:**
 - **Master Database:** Handles write operations and updates.

- **Replicas:** Multiple database replicas are synchronized with the master using replication to handle read-heavy operations and reduce load on the master.
- 5. **Redis:** Used as a caching layer to improve data retrieval speed for authentication publish/subscribe (Pub/Sub)
- 6. **WAL Sync:** Ensures Write-Ahead Logging (WAL) to maintain consistency and reliability across the database replicas.



<https://drive.google.com/file/d/1DtCG5GhppIcJM4azO-Md5Qkdri9ViS1A/view?usp=sharing>

Fig: Distributed Notification and Data Synchronization Architecture

The diagram illustrates a client-server architecture for a notification system, utilizing load balancing, distributed notification replicas, and data synchronization across multiple content database masters. It incorporates:

1. **Client:** Sends requests to the system.
2. **Load Balancer:** Distributes client requests evenly across multiple notification replicas.
3. **Notification Service and Replicas:** Handles event processing and data storage, ensuring redundancy and scalability.
4. **Redis:** Acts as an intermediary for retrieving data in certain scenarios, improving speed.
5. **Content Database Masters:** Store and manage the actual data, with replication mechanisms for fault tolerance.
6. **WAL (Write-Ahead Logging) Sync:** Ensures data consistency and durability across the distributed databases.

Facebook Microservices Feature Collaboration

Communication Technologies Used

1. Synchronous Communication

- **REST APIs:** For direct request-response interactions between services.
- **gRPC:** For high-performance communication, especially for internal service-to-service interactions.
- **GraphQL:** Enables flexible data querying across multiple services, ensuring minimal over-fetching or under-fetching of data.

2. Asynchronous Communication

- **Apache Kafka:** Facilitates event-driven architecture through event streaming and message queuing for decoupled interactions.
- **RabbitMQ:** Provides robust publish/subscribe messaging, particularly for reliable delivery.
- **Redis Pub/Sub:** Powers real-time notifications with low latency.

3. Data Synchronization

- **Change Data Capture (CDC):** Ensures real-time data propagation from one service's database to others.
- **Event Sourcing:** Captures and stores state changes as events to rebuild service states when needed.
- **Write-Ahead Logging (WAL):** Tracks database changes to maintain consistency across replicated databases.

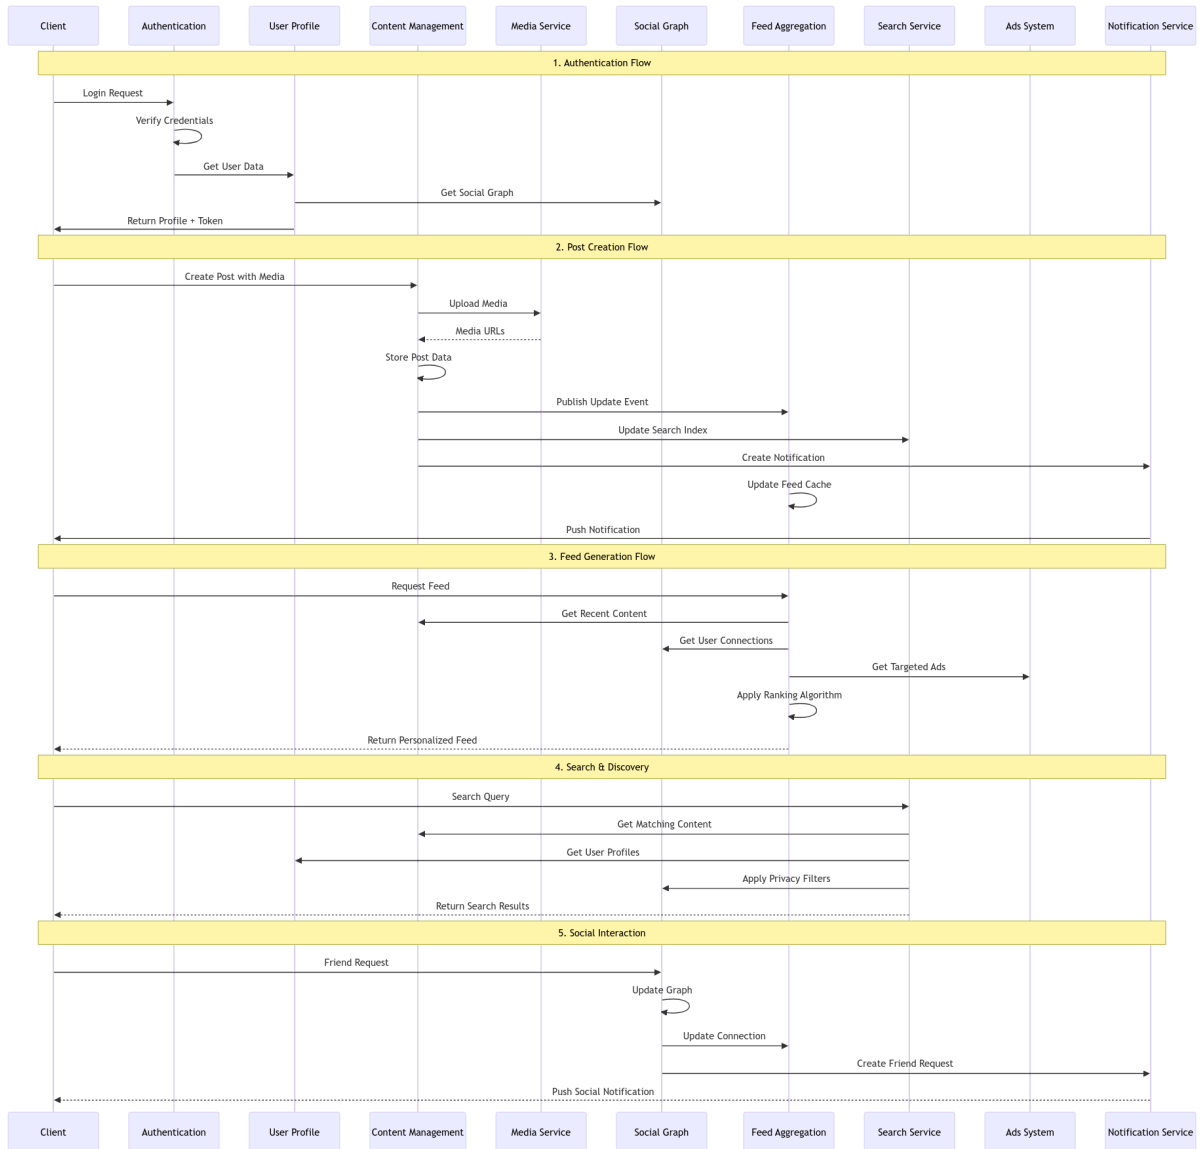


Figure-1 :Complete Facebook Service Interaction Flow

Feature-by-Feature Demonstration

1. User Authentication and Profile Management

Flow:

1. The **Authentication Service** receives a login or registration request.
2. User credentials are validated against the **User Database**.
3. On successful login, the **User Profile Service** loads profile data.
4. The **Privacy Service** applies user-configured privacy settings.
5. A session token (JWT) is generated and cached in Redis for fast validation during subsequent requests.

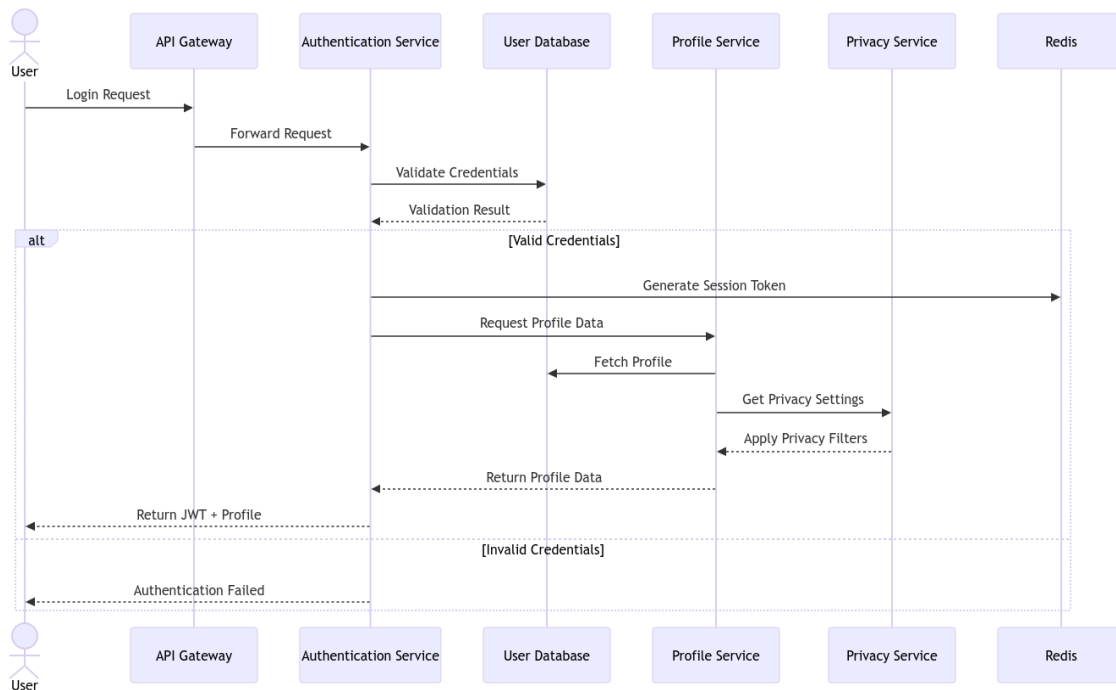


Figure-2 : Authentication and Profile Management Flow

Technologies:

- JWT for secure session management.
- Redis for caching session tokens.
- REST APIs for inter-service communication.
- CDC for syncing changes in the **User Database**.

2. Post Creation and Timeline Update

Flow:

1. **Content Management Service** processes a new post request.
2. Attached media (if any) is uploaded via the **Media Service** and stored in a scalable **Media Database**.
3. The **Post Content Database** is updated with the post details.
4. Events such as **PostCreated** are published to Kafka, triggering:
 - **Feed Aggregation Service**: Updates the user's timeline.
 - **Notification Service**: Notifies followers of the new post.
 - **Search Index DB**: Updates for better discoverability.
5. The **Content Moderation Service** checks the post asynchronously for policy violations.

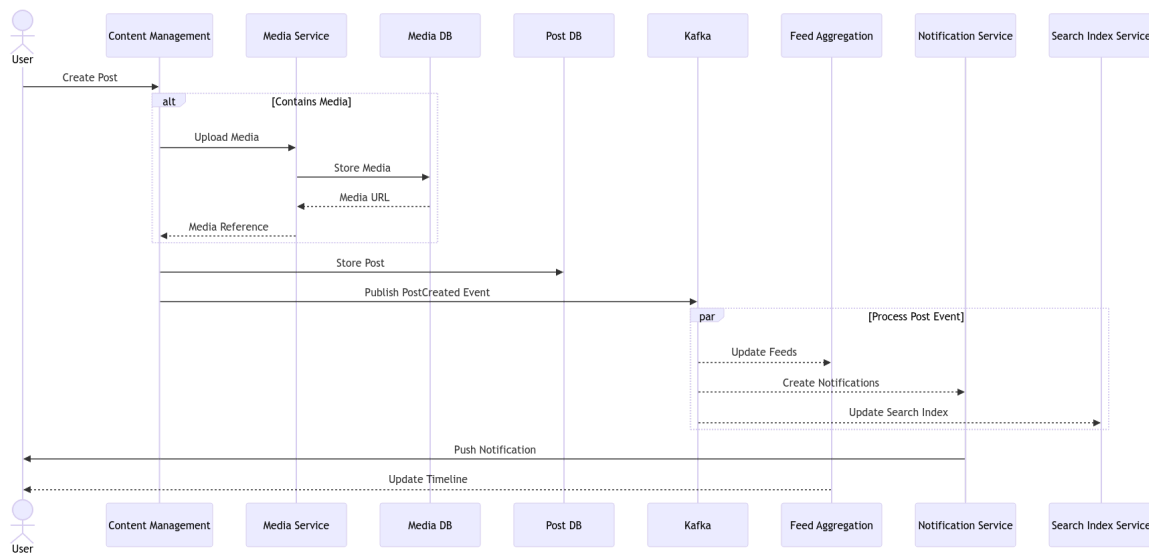


Figure-3 : Post Creation and Timeline Update Flow

Technologies:

- Kafka for event streaming.
- AWS S3 or blob storage for media handling.
- Elasticsearch for search indexing.
- CDC for propagating database updates.

3. Social Graph Operations (Friend Requests/Following)

Flow:

1. A friend request is sent to the **Friendship Service**.
2. The **Friendship Database** records the relationship.
3. Notifications are generated via the **Notification Service**.
4. **Feed Aggregation Service** updates the feed preferences.
5. **Privacy Service** adjusts content visibility settings.

Note: See the next figure of Feed Generation and Social Graph Flow

Technologies:

- Neo4j for managing the social graph.
- Kafka for event propagation.
- Redis for caching frequently accessed friendships.
- REST APIs for service communication.

4. Feed Generation and Content Ranking

Flow:

1. The **Feed Aggregation Service** fetches content from the **Post Content Database** and **Friendship Service**.
2. The **Recommendation Service** personalizes content based on user preferences.
3. The **Ranking Service** orders content by relevance and engagement potential.
4. The **Privacy Service** filters content according to user-defined rules.
5. The finalized feed is cached in the **Feed Cache Database** for efficient delivery.

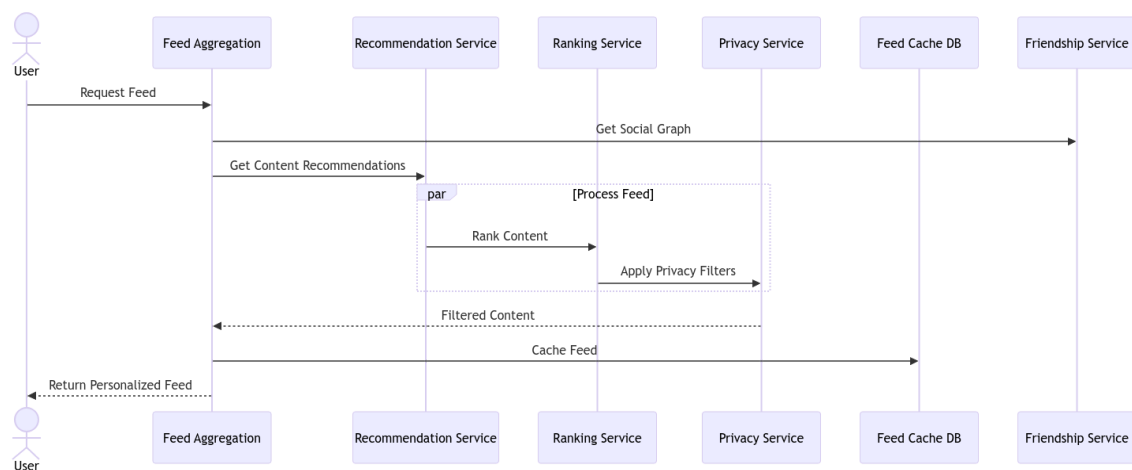


Figure-4 : Feed Generation and Social Graph Flow

Technologies:

- Redis for feed caching.
- gRPC for high-performance data retrieval.
- Machine learning models for personalization and ranking.
- CDC for real-time updates to the feed.

5. Messaging and Real-Time Communication

Flow:

1. The **Messaging Service** receives and processes messages between users.
2. Messages are stored in the **Message Database**.
3. Real-time updates are pushed to the recipient via the **Notification Service**.
4. The **Privacy Service** ensures that message visibility adheres to user settings.

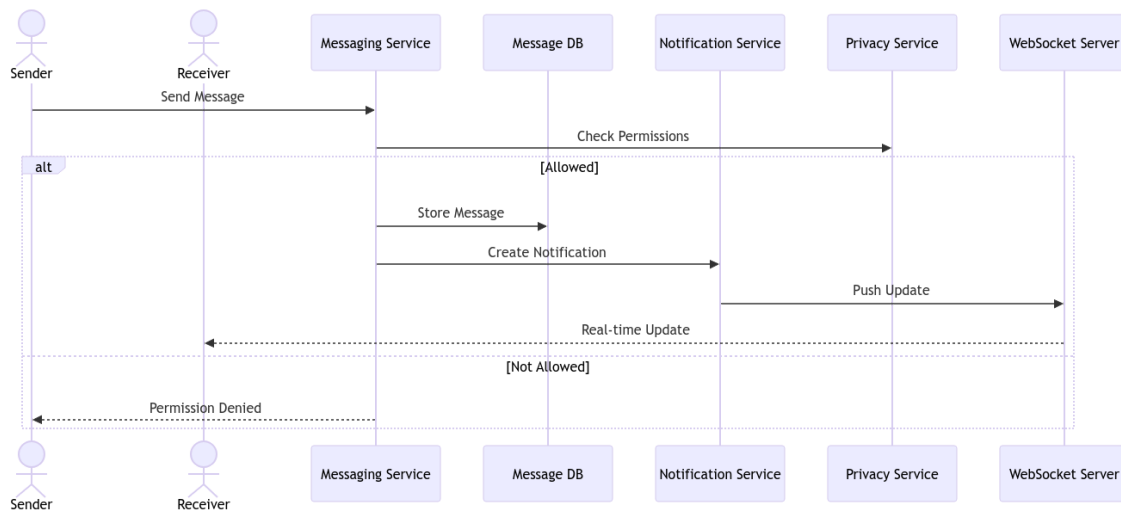


Figure-5 : Messaging and Real-time Communication Flow

Technologies:

- WebSocket for real-time communication.
- Redis Pub/Sub for low-latency notifications.
- RabbitMQ for message queuing and delivery reliability.
- CDC for synchronizing message updates.

6. Ad System Integration

Flow:

1. User activity data is sent to the **Behavior Analysis Service** via Kafka.
2. The **Ads System** processes targeting requests based on behavioral insights.
3. Interactions are logged in the **Event Log**.
4. The **Billing System** manages payment and charges for ad placement.
5. The **Marketplace Service** integrates commercial content into user feeds.
6. Ads are served to users by the **Feed Aggregation Service**.

Note: See the next figure of Ads System and Search Flow

Technologies:

- Kafka for event-driven data pipelines.
- Apache Spark for behavioral analysis.
- PostgreSQL for transactional data.
- Redis for ad targeting and caching.

7. Search Functionality

Flow:

1. The **Search Management Service** processes user queries.
2. It queries the **Search Index Database**, updated periodically via the **Content Management Service**.
3. Results are filtered through the **Privacy Service**.
4. Search results are aggregated from multiple sources, such as user profiles, posts, comments, and marketplace items.

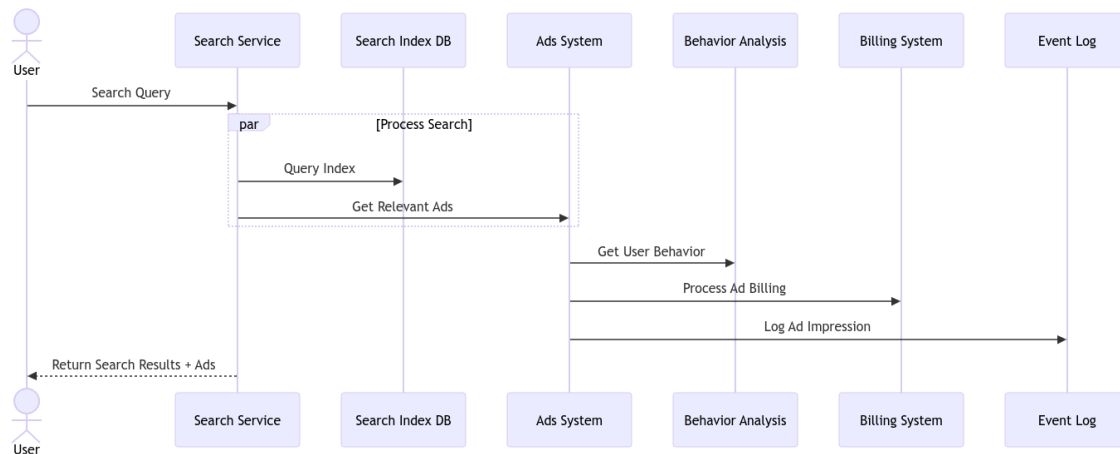


Figure-6 : Ads System and Search Flow

Technologies:

- Elasticsearch for full-text search.
- Redis for caching frequently searched results.
- gRPC for internal communication.
- GraphQL for flexible, aggregated data fetching.

Cross-Cutting Concerns

Data Consistency

- **Eventually Consistent Model:** Used for non-critical updates to ensure scalability.
- **Strong Consistency:** Maintained for sensitive operations like user authentication and financial transactions.
- **CDC:** Ensures real-time synchronization across services.

Security

- **mTLS:** Secures service-to-service communication.
- **API Gateway:** Handles external requests with rate-limiting and validation.
- **JWT:** Authenticates and authorizes users.

Monitoring and Logging

- **Distributed Tracing:** Managed using Jaeger for tracing service requests.
- **Prometheus:** Collects metrics for performance monitoring.
- **ELK Stack:** Aggregates and visualizes logs.

Error Handling

- **Circuit Breakers:** Ensure system resilience by isolating failing services.
- **Retry Mechanisms:** Implemented with exponential backoff.
- **Dead Letter Queues:** Handles undelivered messages gracefully.

Performance Optimization

- Multi-layer caching using Redis.
- Connection pooling for database interactions.
- Load balancing to evenly distribute traffic across service instances.

